

```
#importing the dataset and libraries
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.datasets import mnist
import plotly.graph_objects as go
```

```
#Checking for GPU availability
```

```
if tf.test.gpu_device_name():
    print('GPU device found: {}'.format(tf.test.gpu_device_name()))
    device = '/device:GPU:0'
else:
    print('No GPU device found. Using CPU.')
    device = '/device:CPU:0'
```

```
↳ GPU device found: /device:GPU:0
```

```
(x_train, y_train), (x_test, y_test) = mnist.load_data() #Splitting data into test and train
```

```
x_train = x_train.reshape(-1, 784) / 255.0 #Flattening the (28X28) data
x_test = x_test.reshape(-1, 784) / 255.0 #and normalizing the greyscale
```

```
x_validation, x_train = x_train[:5000], x_train[5000:] #5000 data points for validation
y_validation, y_train = y_train[:5000], y_train[5000:]
```

```
x_train.shape, x_test.shape, y_train.shape, y_test.shape
```

```
((55000, 784), (10000, 784), (55000,), (10000,))
```

```
y_train = tf.keras.utils.to_categorical(y_train)
y_validation = tf.keras.utils.to_categorical(y_validation)
y_test = tf.keras.utils.to_categorical(y_test)
```

```
x_train.shape, x_validation.shape, x_test.shape, y_train.shape, y_validation.shape,
```

```
((55000, 784), (5000, 784), (10000, 784), (55000, 10), (5000, 10), (10000, 10))
```

```

#Neural Network with 3 layers, relu as activation for hidden layers
# with softmax as activation for the output layer
model = Sequential()
model.add(Dense(500, activation='relu', input_shape=(784,)))
model.add(Dense(500, activation='relu'))
model.add(Dense(10, activation='softmax'))

#Compiling the model with Adam Optimizer with categorical cross entropy as the loss
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

#Setting batch size = 128 and epochs to 250
batch_size = 128
epochs = 250

with tf.device(device):

    loss_values = []
    error_values = []

    #each iteration of epochs
    for epoch in range(epochs):

        #each iteration of batch
        for batch_start in range(0, len(x_train), batch_size):
            batch_end = batch_start + batch_size
            x_batch = x_train[batch_start:batch_end]
            y_batch = y_train[batch_start:batch_end]

            model.train_on_batch(x_batch, y_batch)

        #calculating the loss and accuracy
        loss, accuracy = model.evaluate(x_validation, y_validation, batch_size=batch_size)

        #calculating the error
        error = 1 - accuracy
        loss_values.append(loss)
        error_values.append(error)

    print(f"Epoch {epoch+1}/{epochs} - Loss: {loss:.4f} - Error: {error:.4f}")

Epoch 1/250 - Loss: 0.1271 - Error: 0.0398
Epoch 2/250 - Loss: 0.1121 - Error: 0.0358
Epoch 3/250 - Loss: 0.1047 - Error: 0.0324

```

Epoch 4/250 - Loss: 0.0885 - Error: 0.0266  
Epoch 5/250 - Loss: 0.0954 - Error: 0.0244  
Epoch 6/250 - Loss: 0.1207 - Error: 0.0324  
Epoch 7/250 - Loss: 0.0855 - Error: 0.0212  
Epoch 8/250 - Loss: 0.0844 - Error: 0.0210  
Epoch 9/250 - Loss: 0.0972 - Error: 0.0202  
Epoch 10/250 - Loss: 0.0985 - Error: 0.0210  
Epoch 11/250 - Loss: 0.0892 - Error: 0.0188  
Epoch 12/250 - Loss: 0.0855 - Error: 0.0194  
Epoch 13/250 - Loss: 0.1082 - Error: 0.0254  
Epoch 14/250 - Loss: 0.0929 - Error: 0.0210  
Epoch 15/250 - Loss: 0.1037 - Error: 0.0216  
Epoch 16/250 - Loss: 0.1060 - Error: 0.0220  
Epoch 17/250 - Loss: 0.0811 - Error: 0.0178  
Epoch 18/250 - Loss: 0.1030 - Error: 0.0194  
Epoch 19/250 - Loss: 0.1146 - Error: 0.0202  
Epoch 20/250 - Loss: 0.1180 - Error: 0.0190  
Epoch 21/250 - Loss: 0.1154 - Error: 0.0182  
Epoch 22/250 - Loss: 0.1149 - Error: 0.0168  
Epoch 23/250 - Loss: 0.1369 - Error: 0.0226  
Epoch 24/250 - Loss: 0.1372 - Error: 0.0194  
Epoch 25/250 - Loss: 0.1223 - Error: 0.0182  
Epoch 26/250 - Loss: 0.1124 - Error: 0.0172  
Epoch 27/250 - Loss: 0.1212 - Error: 0.0190  
Epoch 28/250 - Loss: 0.1398 - Error: 0.0180  
Epoch 29/250 - Loss: 0.1195 - Error: 0.0178  
Epoch 30/250 - Loss: 0.1180 - Error: 0.0170  
Epoch 31/250 - Loss: 0.1154 - Error: 0.0162  
Epoch 32/250 - Loss: 0.1375 - Error: 0.0196  
Epoch 33/250 - Loss: 0.1263 - Error: 0.0176  
Epoch 34/250 - Loss: 0.1132 - Error: 0.0158  
Epoch 35/250 - Loss: 0.1297 - Error: 0.0178  
Epoch 36/250 - Loss: 0.1437 - Error: 0.0172  
Epoch 37/250 - Loss: 0.1226 - Error: 0.0164  
Epoch 38/250 - Loss: 0.1165 - Error: 0.0154  
Epoch 39/250 - Loss: 0.1341 - Error: 0.0178  
Epoch 40/250 - Loss: 0.1414 - Error: 0.0172  
Epoch 41/250 - Loss: 0.1182 - Error: 0.0148  
Epoch 42/250 - Loss: 0.1548 - Error: 0.0220  
Epoch 43/250 - Loss: 0.1176 - Error: 0.0166  
Epoch 44/250 - Loss: 0.1204 - Error: 0.0162  
Epoch 45/250 - Loss: 0.1636 - Error: 0.0214  
Epoch 46/250 - Loss: 0.1391 - Error: 0.0158  
Epoch 47/250 - Loss: 0.1394 - Error: 0.0152  
Epoch 48/250 - Loss: 0.1361 - Error: 0.0186  
Epoch 49/250 - Loss: 0.1138 - Error: 0.0130  
Epoch 50/250 - Loss: 0.1220 - Error: 0.0148  
Epoch 51/250 - Loss: 0.1517 - Error: 0.0176  
Epoch 52/250 - Loss: 0.1276 - Error: 0.0164  
Epoch 53/250 - Loss: 0.1376 - Error: 0.0150

```
Epoch 54/250 - Loss: 0.1229 - Error: 0.0150
Epoch 55/250 - Loss: 0.1454 - Error: 0.0158
Epoch 56/250 - Loss: 0.1246 - Error: 0.0144
Epoch 57/250 - Loss: 0.1573 - Error: 0.0184
Epoch 58/250 - Loss: 0.1406 - Error: 0.0162
Epoch 59/250 - Loss: 0.1276 - Error: 0.0150
Epoch 60/250 - Loss: 0.1224 - Error: 0.0144
```


# Create the loss plot

```
loss_plot = go.Scatter(x=list(range(1, epochs+1)), y=loss_values, mode='lines', name='Loss')
layout_loss = go.Layout(title='Cross-Entropy Loss', xaxis=dict(title='Epoch'), yaxis=dict(title='Loss'))
fig_loss = go.Figure(data=[loss_plot], layout=layout_loss)
fig_loss.show()
```


# Create the classification error plot

```
error_plot = go.Scatter(x=list(range(1, epochs+1)), y=error_values, mode='lines', name='Error')
layout_error = go.Layout(title='Validation Classification Error', xaxis=dict(title='Epoch'), yaxis=dict(title='Error'))
fig_error = go.Figure(data=[error_plot], layout=layout_error)
fig_error.show()
```

### Cross-Entropy Loss



## Validation Classification Error



[Colab paid products](#) - [Cancel contracts here](#)

---

